

V-TAPP · V-LAUNCHPAD · VIT-AP IIEC & VTBIF

Startup Survivor

Running the event: the day's schedule, how to brief your faculty judges, what every control does, and what to do when something goes wrong.

Organiser Guide

You are the Game Master

6 hours

11:00 – 17:00

40+ teams

3+ judges

What this event actually is

Teams build a startup on a track, then run it through a market simulation four times. Halfway through you fire a shock at everyone, they pivot, and it runs again. Faculty judges grade the documents, and those grades feed back into the simulation.

Your job as Game Master is three things: **keep the clock moving**, **keep the judging queue empty**, and **fire the shock at the right moment**. The system handles fairness on its own.

DESIGN PRINCIPLE WORTH KNOWING

Every mechanism was built so that **no participant can gain from anything except thinking well**.

Clicking fast does nothing. Judges can't favour anyone. Randomness is seeded and reproducible. If a team disputes a result, the full audit log can reconstruct exactly what happened.

Before the day

- 1 **Set up judges.** They're defined in the `.env` file, not in a signup screen. Add a name and code per judge, numbered from 1. However many pairs you add, that many judges exist.
- 2 **Print the judge cards.** Each judge needs their code (e.g. `JUDGE1`) and the URL `/#/judge`. That's their whole login — no password, no account setup.
- 3 **Set the admin passphrase** in `ADMIN_PASSWORD_HASH`. Only you should have it.
- 4 **Run a rehearsal.** `npm run simulate` runs a full 40-team event against a scratch database in about 10 seconds. Do this the week before.
- 5 **Send participants the participant guide** and the Google Doc templates so they arrive prepared.

CHECK THIS ONE SETTING BEFORE THE DOORS OPEN

Every team at the venue shares one WiFi connection, which means they share one public IP address. The login rate limit is set to **300 requests/minute** to handle that. If you lower it, teams will get locked out during registration. Leave `AUTH_RATE_LIMIT` alone unless you know why you're changing it.

The day, hour by hour

11:00 – 11:20 Doors open

Teams register and log in. Judges log in with their codes. Chase anyone whose pre-event documents are missing.

11:20 – 12:05 Tick 1

Idea & Problem Statement due. 20 min to submit, 25 min for judges, then the tick fires.

12:05 – 12:50 Tick 2

Pitch Deck due — the last thing before the shock.

12:50 – 13:05 ⚡ Halftime shock + the cut

You fire this. It locks pre-shock documents and disqualifies teams that submitted nothing.

13:05 – 13:50 Pivot window

45 minutes. Teams submit their Pivot Rationale — the biggest single score in the event.

13:50 – 14:35 Tick 3

Financial Model and Go-to-Market Plan due.

14:35 – 15:20 Tick 4

Final Pitch Deck due.

15:20 – 16:00 Judging catch-up

Judges finish anything outstanding. No new ticks.

16:00 – 16:40 Freeze results

Lock the scoring weights, verify the standings, handle any disputes.

16:40 – 17:00 🏆 Results

Reveal the standings and hand out prizes.

EACH 45-MINUTE TICK, INTERNALLY

20 min teams submit → **25 min** judges grade → **tick fires**. The tick will not run until grading is finished, so the 25 minutes is a real deadline for your judges, not a suggestion.

Briefing your faculty judges

Faculty are not gamers and shouldn't need to understand the simulation. Brief them on four things and nothing else.

1. Logging in takes five seconds

Go to `/#/judge`, type your code, done. No password, no account. Each judge has their own code so their scores are attributed to them.

2. There are two completely different jobs

JOB	WHAT YOU DO	WHY
Pricing Justification	Approve or reject. No score.	This decides whether the team's new numbers apply this tick. It's a validity check: do the citations support the price?
Everything else	Award points. No approve/reject.	These are graded for quality and feed the team's score.

The system enforces this — a judge physically cannot reject a pitch deck or put a score on a pricing doc.

3. Deadlines are real

Judges have **25 minutes per tick**. The market will not run until every pricing decision is made and every document is rated. If judges fall behind, the whole event stalls. Tell them this explicitly.

4. What they see when reviewing

Before the document, the screen shows how that team is actually doing — money made or lost, revenue, points so far, and an insolvency warning if relevant. They're never rating a pricing doc blind.

THINGS YOU DON'T NEED TO EXPLAIN TO FACULTY

They can't break anything. Judges have no access to start/stop the game, reset anything, fire targeted shocks, or change scoring. Two judges opening the same submission is handled automatically — the second one sees it's claimed and moves on. Multiple judges scoring the same team is expected and averaged.

Judge allocation happens automatically

TEAMS	WHAT HAPPENS
Under 20	Every judge sees every team. Scores are averaged.
20 or more	Teams are split evenly between judges (40 teams / 3 judges = 14/13/13). Because each team is then seen by one judge, scores are statistically normalised so a harsh judge doesn't unfairly sink their teams.

The mode locks at the first tick so it can't change mid-event. Late-registering teams are auto-assigned to whichever judge has the fewest.

Your controls

Everything below is at `/#/gm`, admin only.

CONTROL	WHAT IT DOES	WHEN
Start / Pause / Resume / End	Game phase	Start at 11:20, end after tick 4
Run tick	Fires the market for everyone	End of each tick window
Force tick	Runs the tick despite ungraded work	EMERGENCY ONLY
Fire halftime shock	The shock + document lock + the cut, all at once	12:50, once
Extend deadlines	Pushes every deadline for the current window	When running late
Advance round	Moves to round 2	After the pivot window
Freeze results	Locks scoring weights permanently	Before the reveal
Reset	Wipes all teams	NEVER, DURING AN EVENT

THE TICK WON'T RUN — THIS IS NORMAL

If you press Run tick and get **"Grading is not finished"**, that's the system working correctly. Document points feed the market, so running early would compute results from a half-graded field.

The panel lists exactly what's outstanding — which team, which document, and whether it needs a decision or a rating. Go tell those judges. Only use **Force tick** if a judge has genuinely gone missing; it's logged permanently with a list of what was skipped.

Firing the halftime shock

This is the centrepiece of the event and it does four things in one click, so do it deliberately.

- 1 **Preview the cut first.** The panel shows exactly who would be disqualified before you commit to anything.
- 2 **Check the "at risk" list.** These are teams below the bar only because judges haven't graded them yet. They are **never** auto-cut — chase the grading first if you want them assessed fairly.
- 3 **Fire it.** The shock hits every team equally, pre-shock documents lock permanently, and the cut executes.
- 4 **Announce it in the room.** This is the dramatic moment of the day — don't let it happen silently on 40 laptops.

THE CUT RULE

Teams that submitted **zero documents** before halftime are disqualified. It's a participation gate, not a quality bar — anyone who did the work stays in. Disqualified teams keep their login but can't submit or deploy, and rank below all surviving teams.

You can **reinstate** any team if the cut caught someone unfairly.

When something goes wrong

PROBLEM	WHAT TO DO
Judges are behind and the tick is blocked	Extend deadlines to buy time. If a judge has genuinely vanished, Force tick .
A team's laptop died	They lose nothing — an unsubmitted tick carries over their last approved strategy automatically.
A team says their document is rejected unfairly	You can override any judge decision. It's logged with your name.
A team was cut but shouldn't have been	Reinstate them from the admin panel.
A team disputes their final placing	Open <code>/#/logs</code> , filter by their team, and walk through it. Every action is timestamped.
A team says "we made the most money, why aren't we first?"	Money is half the score. Show them the standings breakdown — it lists money, revenue, documents, judge score and market score side by side.
Running badly behind schedule	Reduce the grading window, or drop tick 4. The schedule is configurable.

The logs

At `/#/logs`, open to you and your judges. Two tabs:

- **Event log** — every action ever taken, filterable by team or type, exportable as CSV. A 40-team event generates around 500 entries.
- **Final standings** — the full scoring breakdown per team.

Entries read in plain English, e.g. "*Team08 disqualified at halftime: Submitted 0 of the 1 document(s) required*" or "*Tick 1 carried over team 3's previous strategy (no approved request)*". This is your dispute-resolution tool — export the CSV at the end of the day and keep it.

How scoring works (for when you're asked)

Two halves, each out of 100, weighted evenly by default:

50%

MARKET PERFORMANCE

50%

DOCUMENT POINTS

Market performance

Each tick, every team is scored **0–100 against the other teams on their track**, based on what they actually kept (revenue minus costs). Those four scores are then averaged.

This matters: with only four ticks, using raw rupees would let a single lucky tick decide the whole event. Averaging normalised scores means all four count equally. It also means a team on a tough track isn't punished for their track's baseline — they're only ever compared to their direct rivals.

Document points

Judge-awarded, totalling exactly 100 across the six graded documents. These also feed back into the simulation as up to a $\pm 30\%$ swing on how attractive a team is to customers — which is how "judges affect the outcome" is actually implemented.

What's deliberately not counted twice

The 1–10 rubric scores judges give already influence the market through customer retention. They aren't added again as a third scoring component, because that would be counting the same opinion twice.

WEIGHTS ARE ADJUSTABLE — BUT FREEZE THEM

You can change the 50/50 split before results are revealed. Once you press **Freeze results** they're locked permanently. Freeze before you look at the standings, so nobody can suggest the weights were tuned to pick a winner.

Fairness guarantees you can state publicly

- **Clicking fast gains nothing.** All teams tick simultaneously.
- **Judges cannot favour anyone.** Events and powers hit the entire field; there's no targeting option in their interface at all.
- **Results are reproducible.** Randomness is seeded, so the same inputs always produce the same outputs.
- **Judge timing doesn't decide winners.** The tick waits for grading, so nobody benefits from being graded early.
- **No team is punished for a slow queue.** Ungraded work never causes a cut or a penalty.

- **Teams can't see each other's standings** until you release them.
- **Everything is logged** in an append-only trail written in the same transaction as the action, so the log can never disagree with the data.

Quick reference

WHO	URL	LOGIN
Participants	/	Register / team password
Faculty judges	/#/judge	Their code (JUDGE1, JUDGE2...)
You	/#/gm	Admin passphrase
Logs	/#/logs	Staff only

The judge and Game Master doors are **not linked from the landing page**. Participants only see register and team login. Share those URLs with staff directly.

YOUR FIVE JOBS ON THE DAY

1. Start the game at 11:20. **2.** Run each tick when grading clears. **3.** Fire the halftime shock at 12:50 and announce it loudly. **4.** Keep the judging queue empty — chase judges, extend if needed. **5.** Freeze results before you look at them, then reveal.